

Swift Key :Capstone : Data Cleaning and Processing Week1

Kumar Shaket

20/06/2021

Week 1

Loading Natural Language Processing and ggplot packages

##Loading text files after reading file

```
blogfile <- "/Users/kumarshaket/Desktop/Coursera/Capstone Project/final/en_US/en_US.blogs.txt"
newsfile <- "/Users/kumarshaket/Desktop/Coursera/Capstone Project/final/en_US/en_US.news.txt"
twitterfile <- "/Users/kumarshaket/Desktop/Coursera/Capstone Project/final/en_US/en_US.twitter.txt"
bloglines <- readLines(blogfile,encoding = "UTF-8",skipNul = TRUE)
newslines <- readLines(newsfile,encoding = "UTF-8",skipNul = TRUE)
twitterlines <- readLines(twitterfile,encoding = "UTF-8",skipNul = TRUE)
```

2.The en_US.twitter.txt has how many lines of text? Is Over 2 million

```
blogword_count <- length(bloglines)
blogword_count
```

```
## [1] 899288
```

```
twitterword_count <- length(twitterlines)
twitterword_count
```

```
## [1] 2360148
```

```
newsword_count <- length(newslines)
newsword_count
```

```
## [1] 1010242
```

3.What is the length of the longest line seen in any of the three en_US data sets? Is Over 40 thousand in the blogs data set

```
max(nchar(newslines))
```

```
## [1] 11384
```

```
max(nchar(bloglines))
```

```
## [1] 40833
```

4. In the en_US twitter data set, if you divide the number of lines where the word “love” (all lowercase) occurs by the number of lines the word “hate” (all lowercase) occurs, about what do you get?

```
love_count <- sum(grep("love",twitterlines))  
hate_count <- sum(grep("hate",twitterlines))  
love_count/hate_count
```

```
## [1] 4.111791
```

5. The one tweet in the en_US twitter data set that matches the word “biostats” says what? As below

```
biostat <- grep("biostats",twitterlines)  
twitterlines[biostat]
```

```
## [1] "i know how you feel.. i have biostats on tuesday and i have yet to study =/"
```

##6. How many tweets have the exact characters “A computer once beat me at chess, but it was no match for me at kickboxing”. (I.e. the line matches those characters exactly.) Is 3

```
sentenceTwitter <- grep("A computer once beat me at chess, but it was no match for me at kickboxing",twitterlines)  
length(sentenceTwitter)
```

```
## [1] 3
```

Week 2

Sampling the Data

```
set.seed(3000)  
stwitter <- sample(twitterlines,size = 2000,replace= TRUE)  
sblog <- sample(bloglines,size=2000,replace = TRUE)  
snews <- sample(newslines,size = 2000,replace = TRUE)
```

A corpus is created using above sample data from twitter, blog and news

```
corpus <- VCorpus(VectorSource(c(stwitter,snews,sblog)),readerControl = list(readPlain,language="en",lo
```

Exploratory and Cleaning of Data

This section will use the text mining library ‘tm’ (loaded previously) to perform Data cleaning tasks, which are meaningful in Predictive Text Analytics. Main cleaning steps are:

1. Converting the document to lowercase.
2. Removing Whitespace.
3. Removing Punctuation.
4. Removing Numbers 5. Removing stopwords (i.e. “and”, “or”, “not”, “is”)
5. Removing Profanity words from data The above can be achieved with some of the TM package functions; let’s take a look to each cleaning task, individually:
6. Reading profanity words list downloaded from internet and storing in variable

```
badwords <- readLines("/Users/kumarshaket/Desktop/Coursera/Capstone Project/Coursera-Capstone-Project/b
```

```
## Converting to lower case
corpus_lower <- tm_map(corpus,content_transformer(tolower))
## Removing whitespace from data
corpus_space <- tm_map(corpus_lower,stripWhitespace)
## Removing Punctuation from data
corpus_punc <- tm_map(corpus_space,removePunctuation)
## Removing Number from data
corpus_num <- tm_map(corpus_punc,removeNumbers)
## Removing stop words from data
corpus_nostop <- tm_map(corpus_num,removeWords,stopwords("english"))
## Removing profanity words from data
finalCorpus <- tm_map(corpus_nostop,removeWords,badwords)
```

Tokenization

Let’s read the text to break it into words and sentences, and to turn it into n-grams. These are all called tokenization because we are breaking up the text into units of meaning, called tokens.

In Natural Language Processing (NLP), n-gram is a contiguous sequence of n items from a given sequence of text or speech. Unigrams are single words. Bigrams are two words combinations. Trigrams are three-word combinations.

The tokenizer method is allowed in R using the package RWeka. The following function is used to extract unigram, bigram, trigram from the text Corpus using RWeka.

Obtaining Unigram

```

unigram <- as.data.frame((as.matrix(TermDocumentMatrix(finalCorpus))))
unigram1v <- sort(rowSums(unigram),decreasing = TRUE)
unigram1d <- data.frame(word=names(unigram1v),freq=unigram1v)
unigram1d[1:10,]

```

```

##      word freq
## said said  571
## will will  506
## one  one  476
## just just  457
## like like  430
## can  can  394
## time time  375
## get  get  355
## new  new  286
## good good  264

```

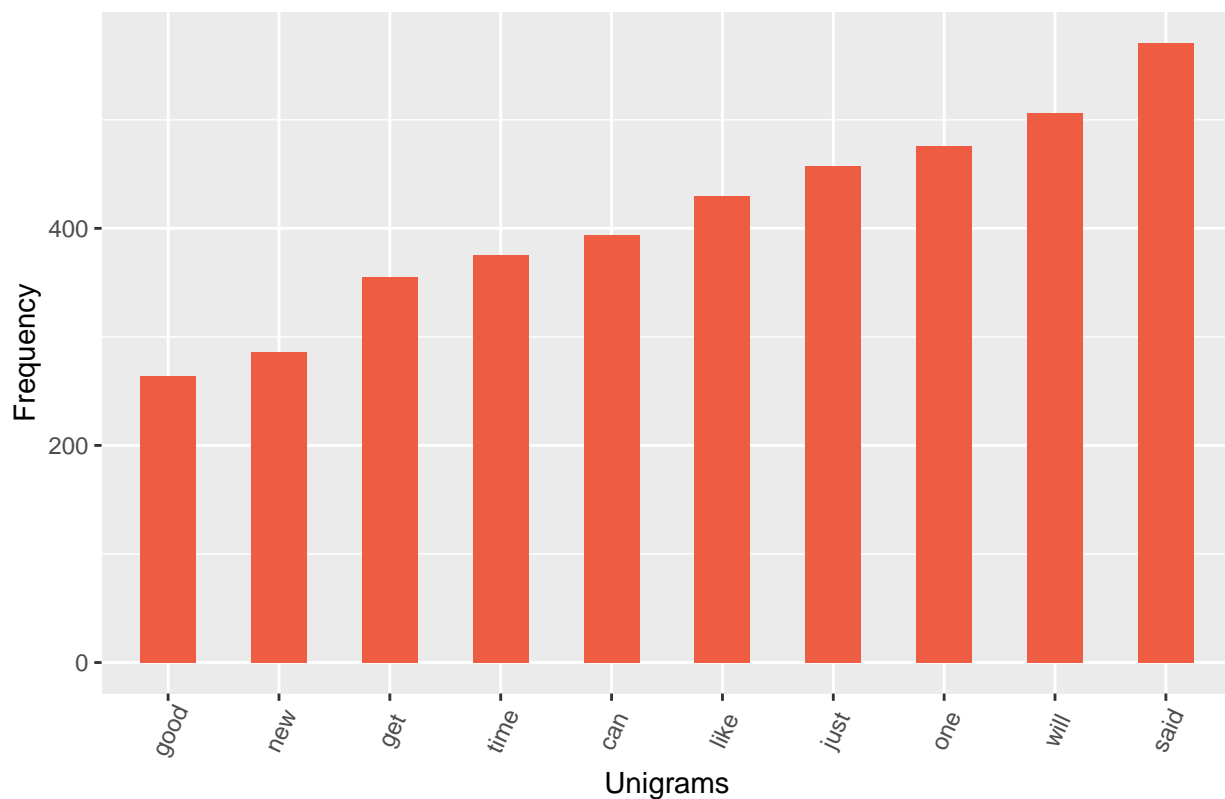
Histogram of top 10 Unigrams

```

ggplot(unigram1d[1:10,],aes(x=reorder(word,freq),y=freq))+
  geom_bar(stat = "identity",width = 0.5,fill="tomato2") +
  labs(title = "Unigrams")+
  xlab("Unigrams")+ylab("Frequency")+
  theme(axis.text.x = element_text(angle = 65,vjust = 0.6))

```

Unigrams



Obtaining BiGram


```
bigram <- function(x) NGramTokenizer(x,Weka_control(min=2,max=2))
gram2 <- as.data.frame(as.matrix(TermDocumentMatrix(finalCorpus,control = list(tokenize=bigram))))
gram2v <- sort(rowSums(gram2),decreasing = TRUE)
gram2d <- data.frame(word=names(gram2v),freq=gram2v)
gram2d[1:10,]
```

```
##              word freq
## right now    right now  37
## last year    last year  35
## last week    last week  28
## years ago    years ago  27
## first time   first time  24
## st louis     st louis   24
## cant wait    cant wait  22
## make sure    make sure  22
## dont know    dont know  21
## last night   last night  21
```

Histogram of top 10 BiGram

```
ggplot(gram2d[1:10,],aes(x=reorder(word,freq),y=freq))+
  geom_bar(stat = "identity",width = 0.5,fill="blue") +
  labs(title = "Bigrams")+
  xlab("Bigrams")+ylab("Frequency")+
  theme(axis.text.x = element_text(angle = 65,vjust = 0.6))
```

Bigrams



Obtaining Trigram

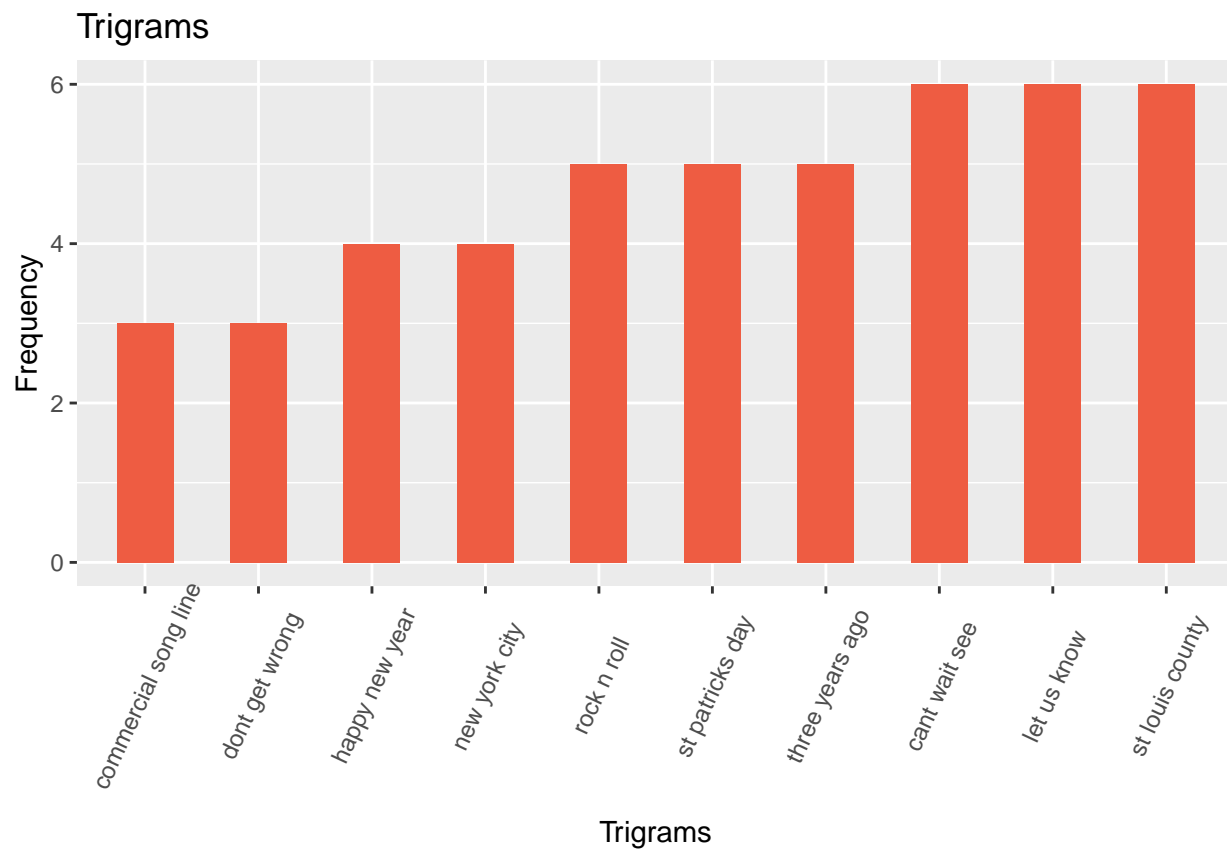
```
trigram <- function(x) NGramTokenizer(x, Weka_control(min=3, max=3))
gram3 <- as.data.frame(as.matrix(TermDocumentMatrix(finalCorpus, control = list(tokenize=trigram))))
gram3v <- sort(rowSums(gram3), decreasing = TRUE)
gram3d <- data.frame(word=names(gram3v), freq=gram3v)
gram3d[1:10,]
```

```
##                                word freq
## cant wait see                 cant wait see    6
## let us know                   let us know      6
## st louis county               st louis county   6
## rock n roll                   rock n roll      5
## st patricks day               st patricks day   5
## three years ago               three years ago   5
## happy new year                happy new year    4
## new york city                 new york city     4
## commercial song line          commercial song line 3
## dont get wrong                dont get wrong     3
```

Histogram Of 10 Top Trigrams

```
ggplot(gram3d[1:10,], aes(x=reorder(word, freq), y=freq)) +
  geom_bar(stat = "identity", width = 0.5, fill="tomato2") +
  labs(title = "Trigrams") +
```

```
xlab("Trigrams")+ylab("Frequency")+
  theme(axis.text.x = element_text(angle = 65,vjust = 0.6))
```



Next Steps

1. Build a Shiny app to allow the user input the word to obtain a suggestion of the next word.
2. Develop the prediction algorithm implemented in Shiny app.
3. Prepare a pitch about the app and publish it at “shinyapps.io” server.